

Introducing Pub Sub Systems

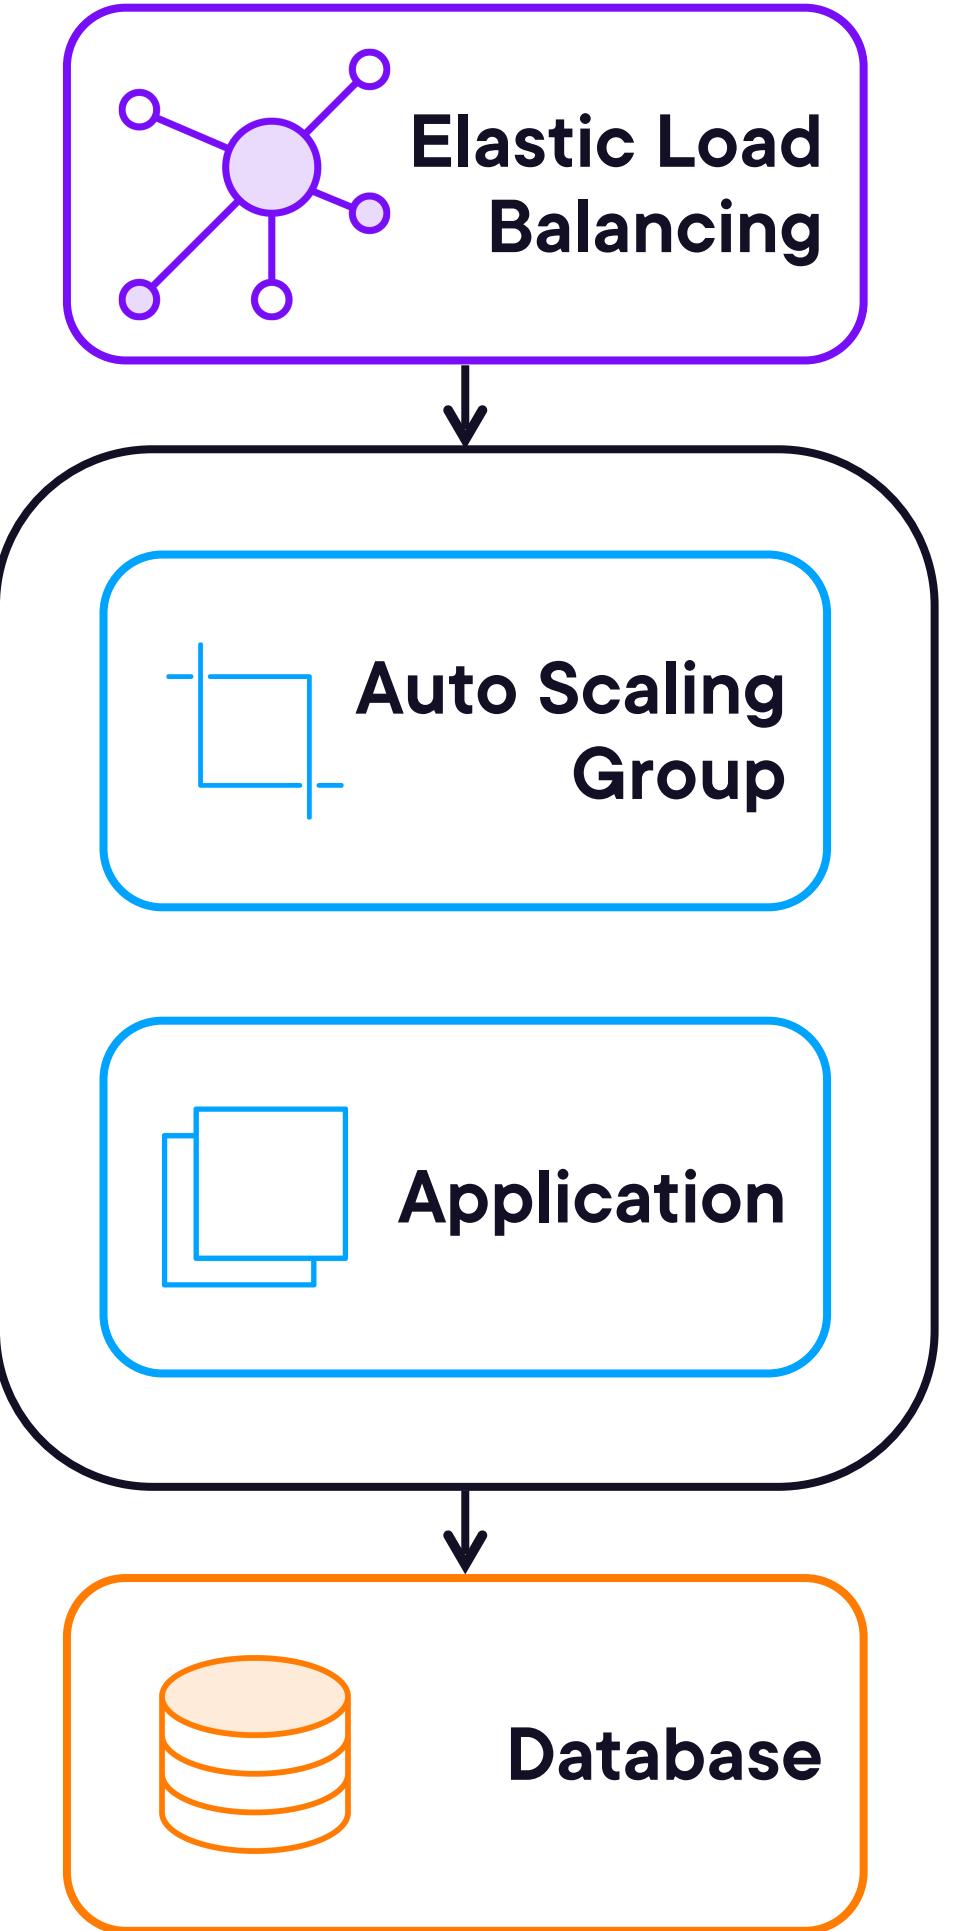
Meeting Pub-Sub Systems

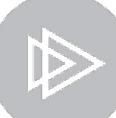
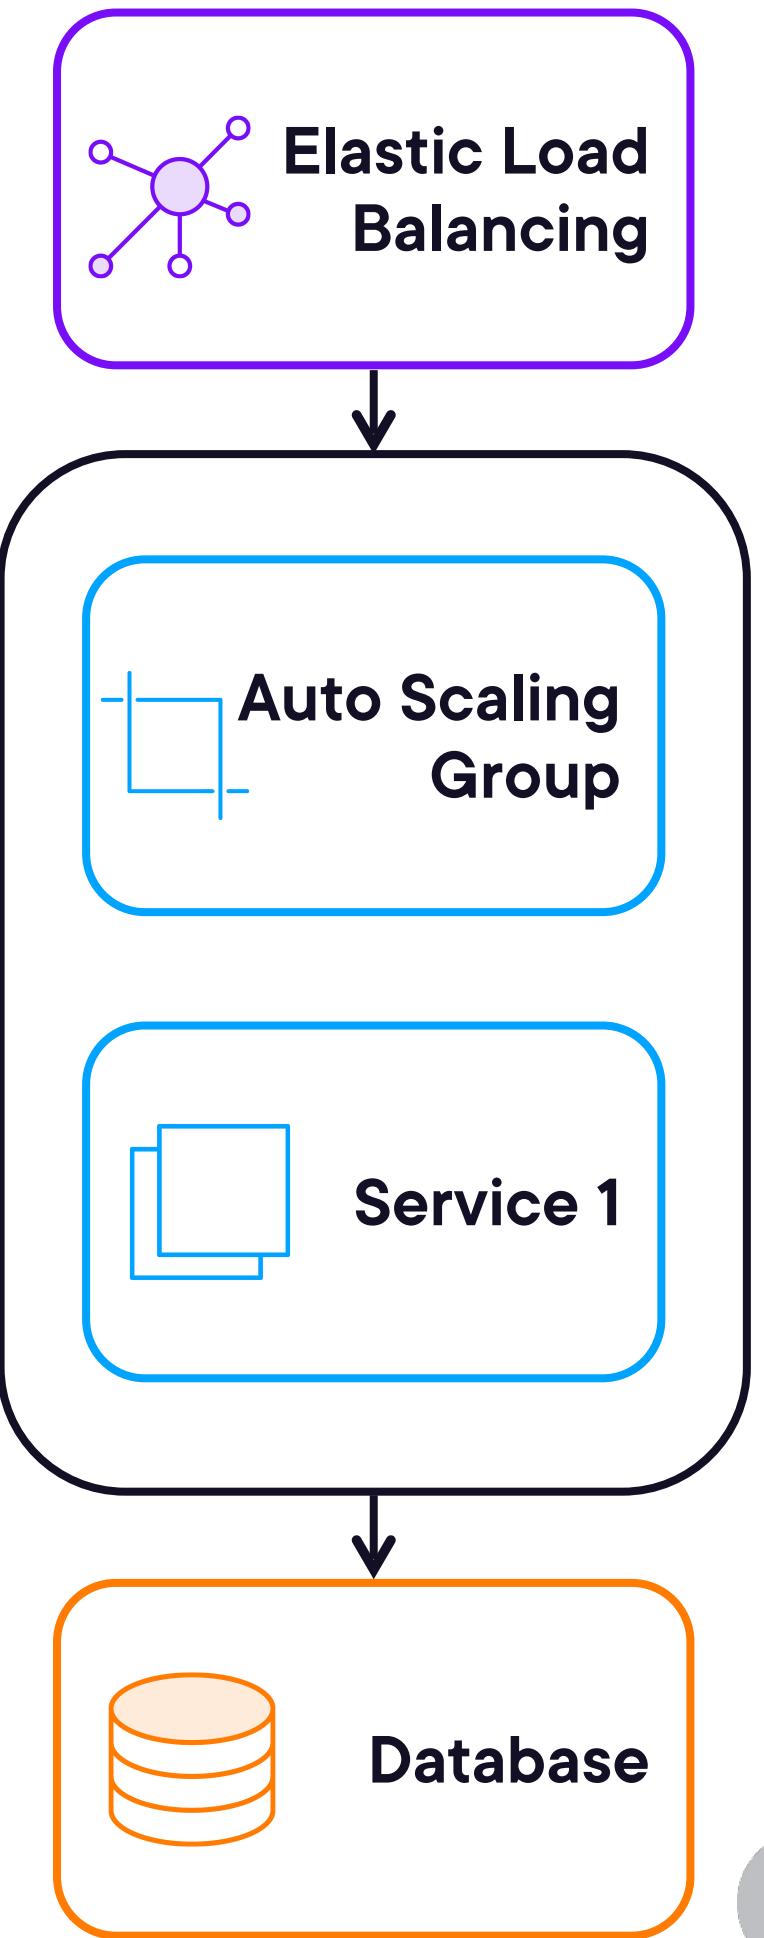
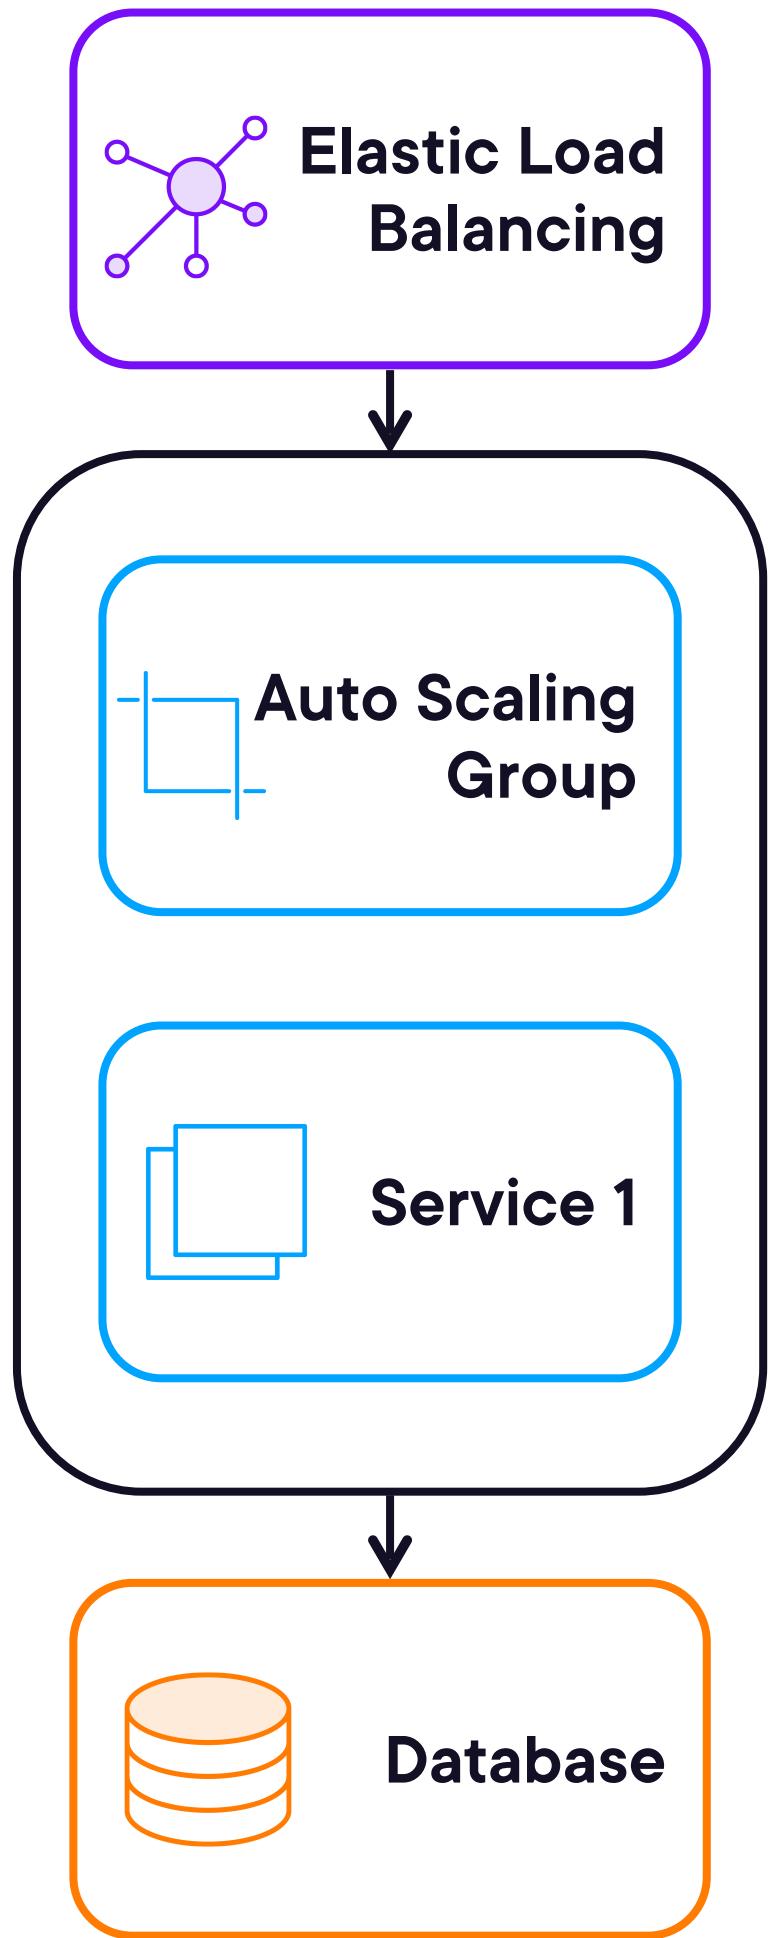
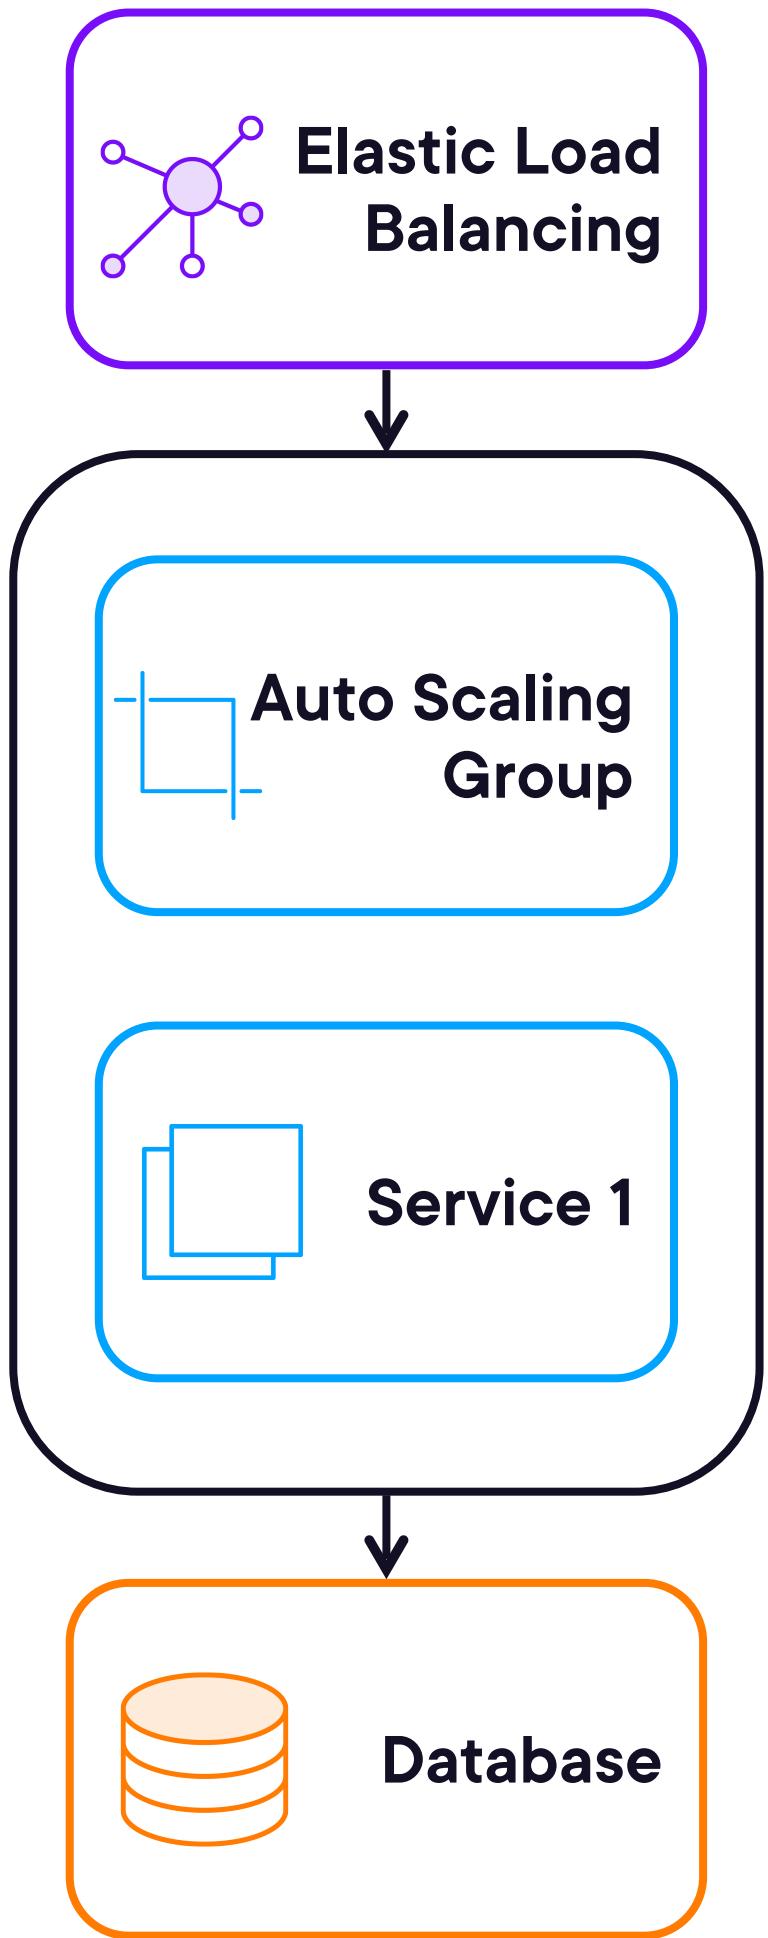
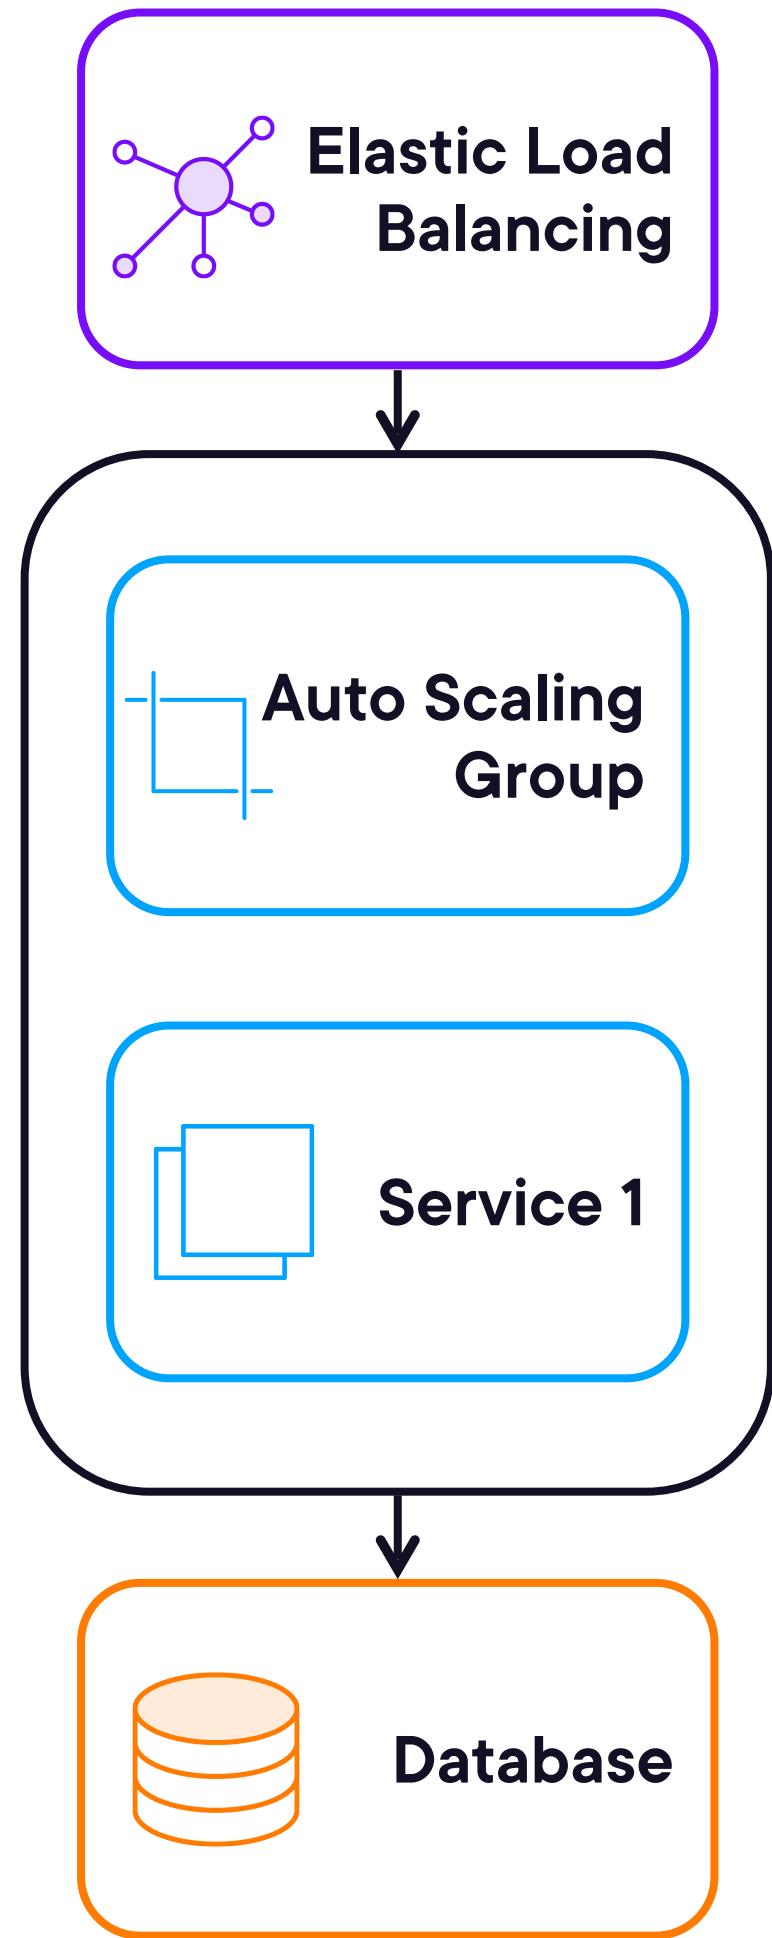


Axel Sirota

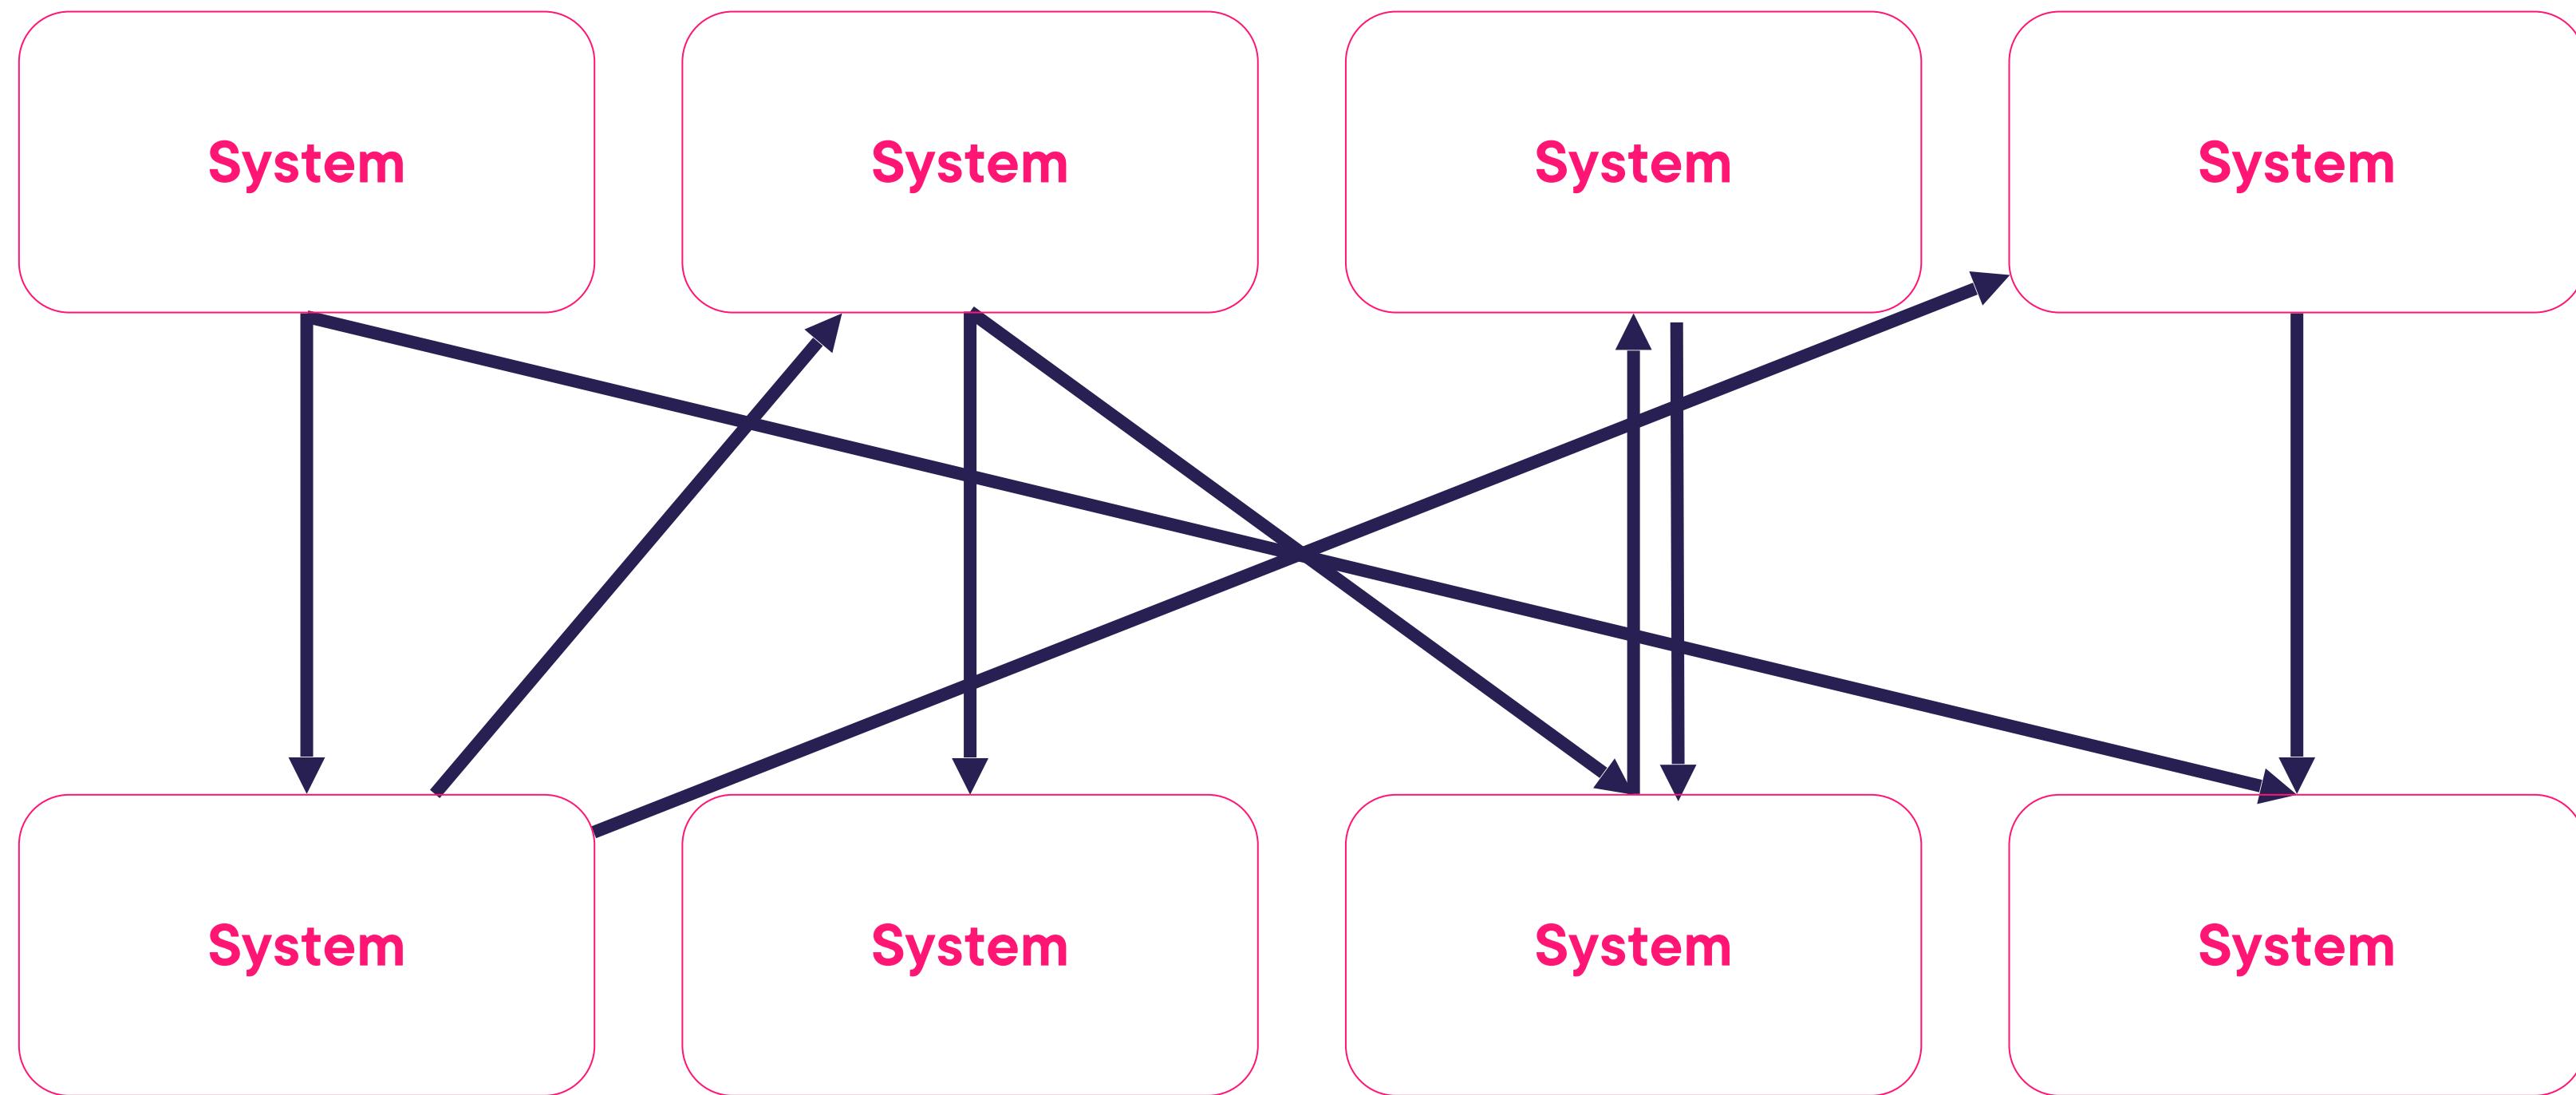
AI and Cloud Consultant

@AxelSirota





Entangled System

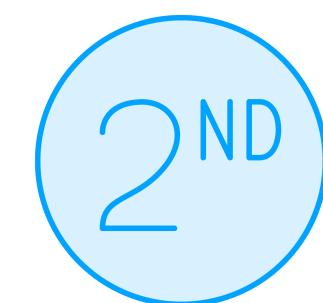




We Call:



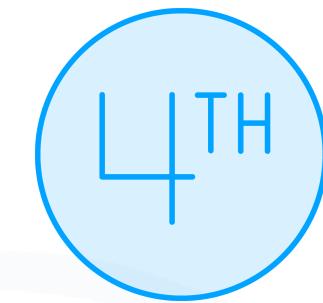
The entity/app creating a message, a publisher or producer



The entity/app consuming messages from a channel as a consumer



The system where the channels live and handle this requests as an Event Bus or, more recently, streaming platforms



The channel where messages flow as channel or topic



Back to Definitions

1ST

We say the pub-sub system is reliable when you ensure there is no message loss

2ND

Has at most once processing when you ensure there is no message duplication

3RD

And has exactly-once processing when you only process a message once ensuring it wasn't lost. Of course, the holy grail.



We Had Ton of Other Pub/Sub Systems Before:

RabbitMQ

Mulesoft

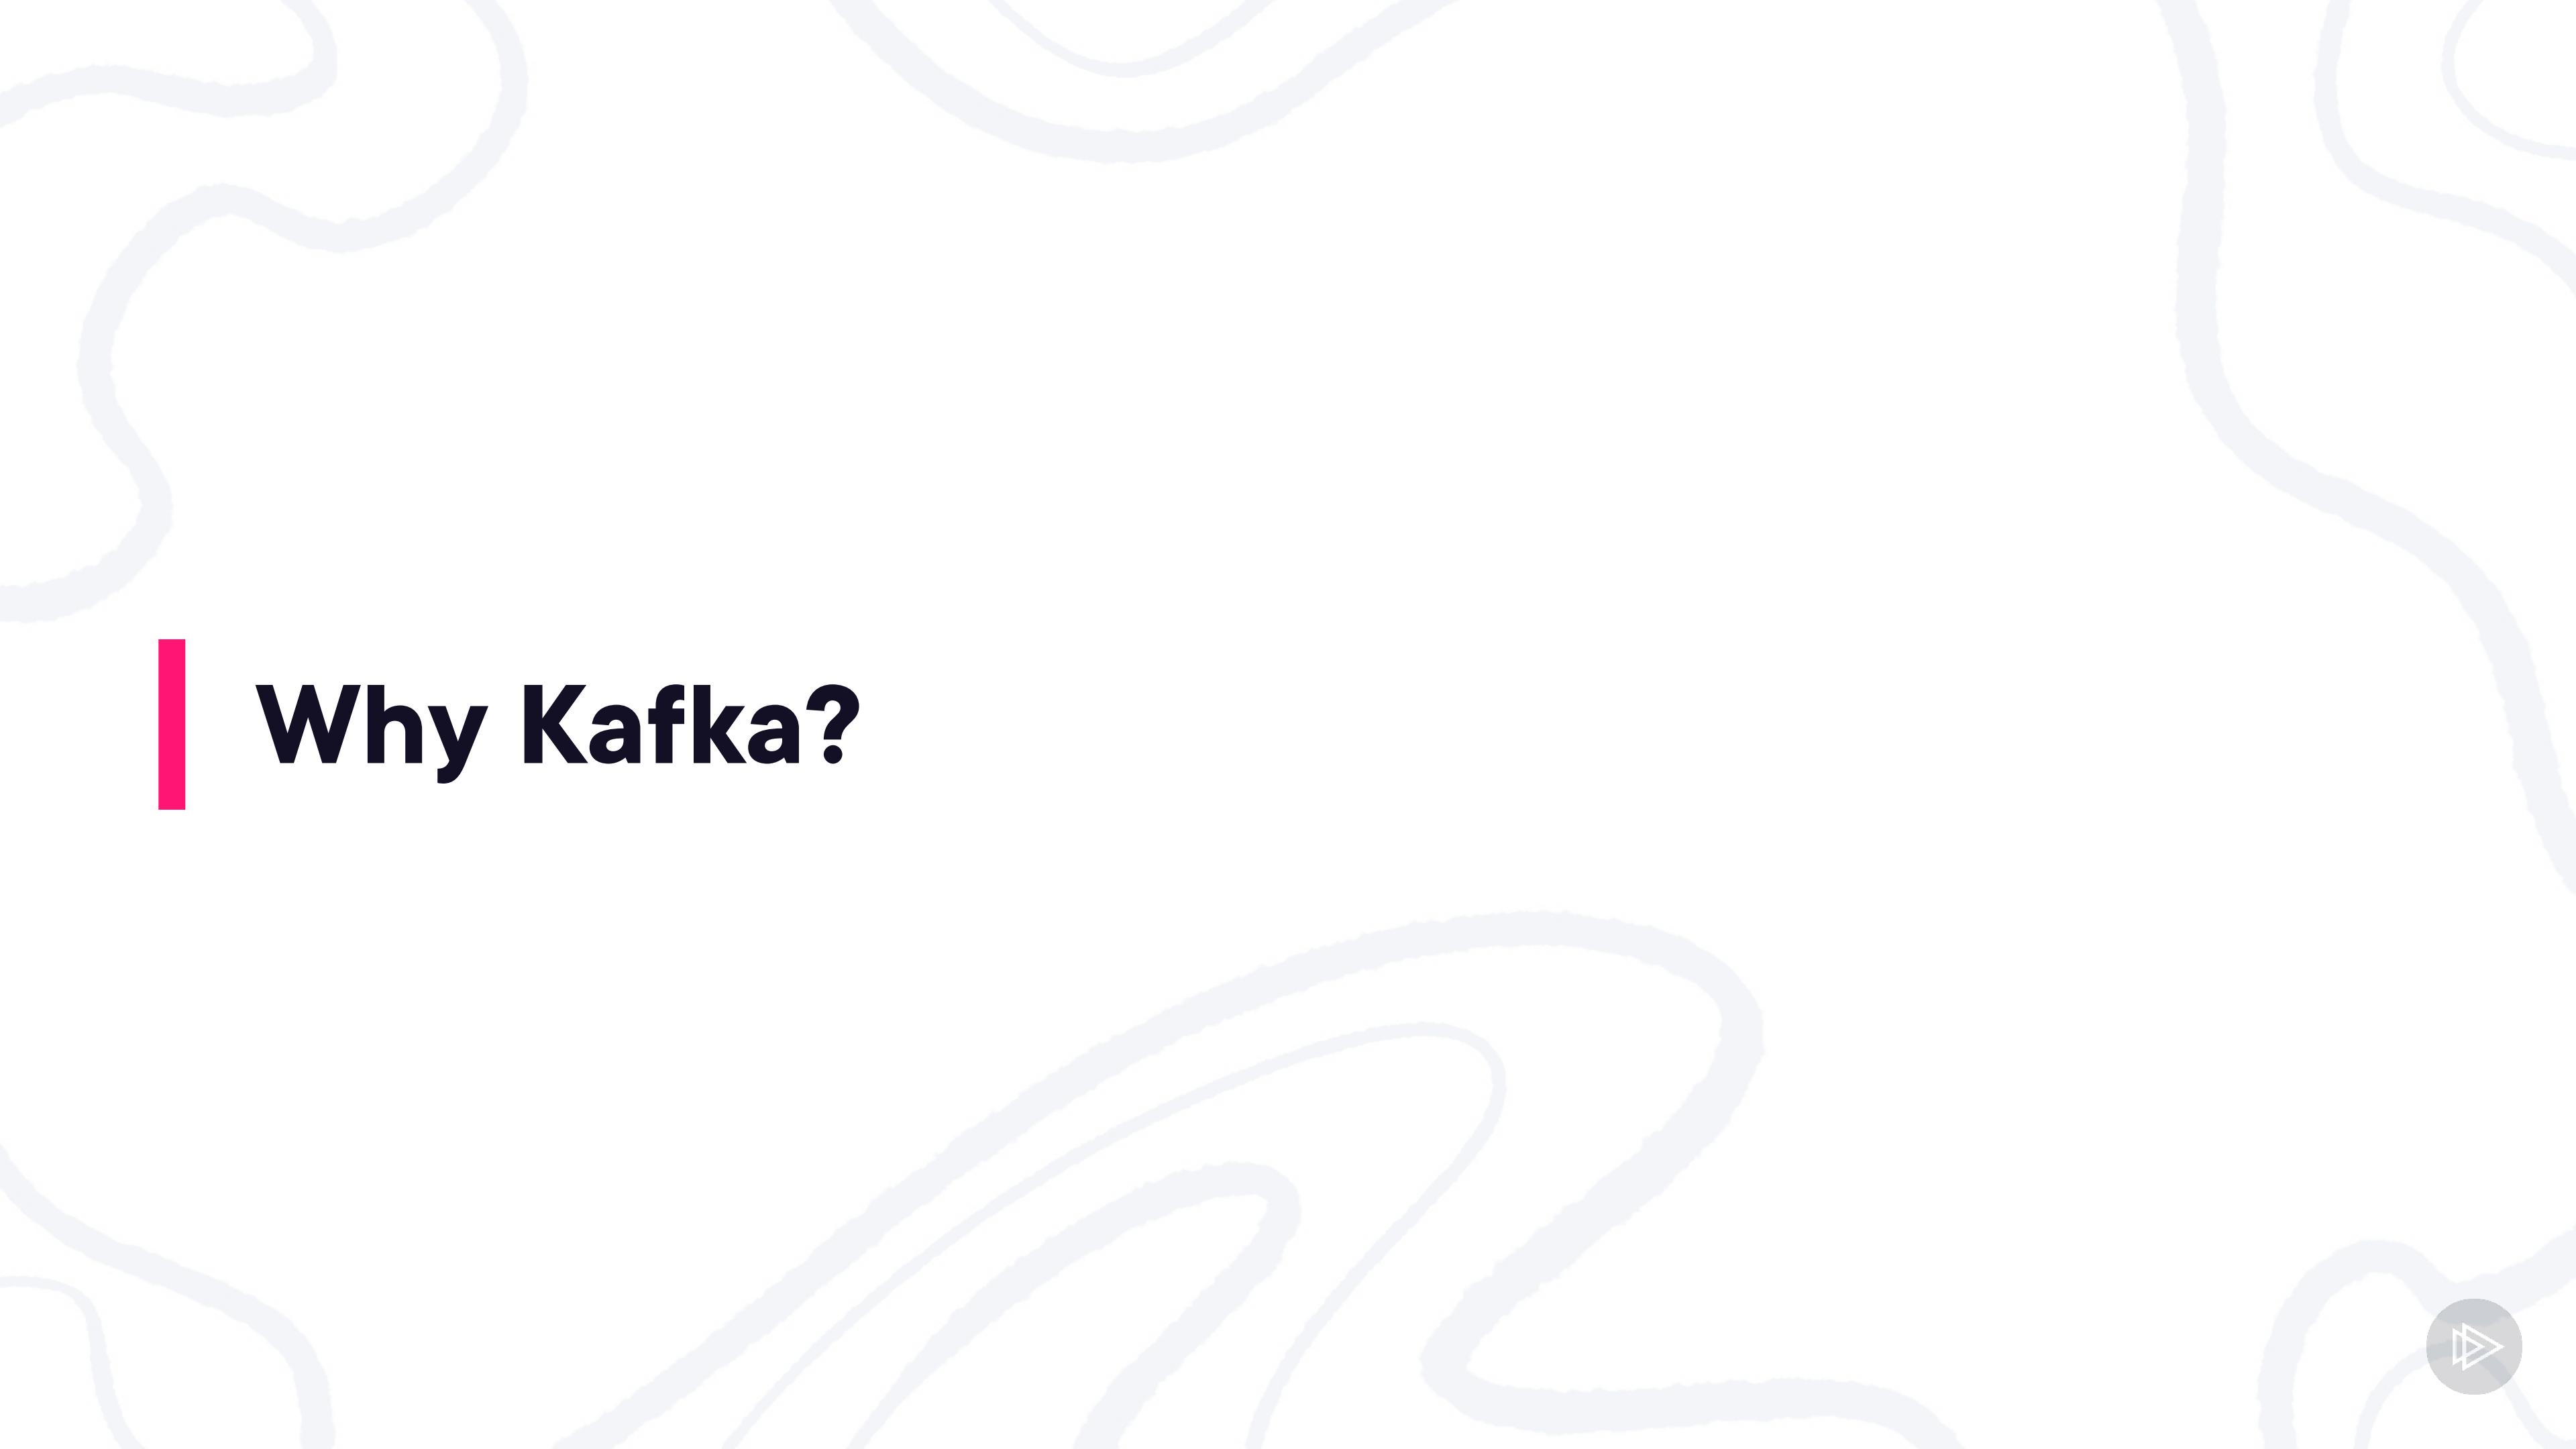
Redis

Amazon SQS

Amazon SNS

Azure HDInsight, Azure
Streaming Service, and
Google Pub/Sub





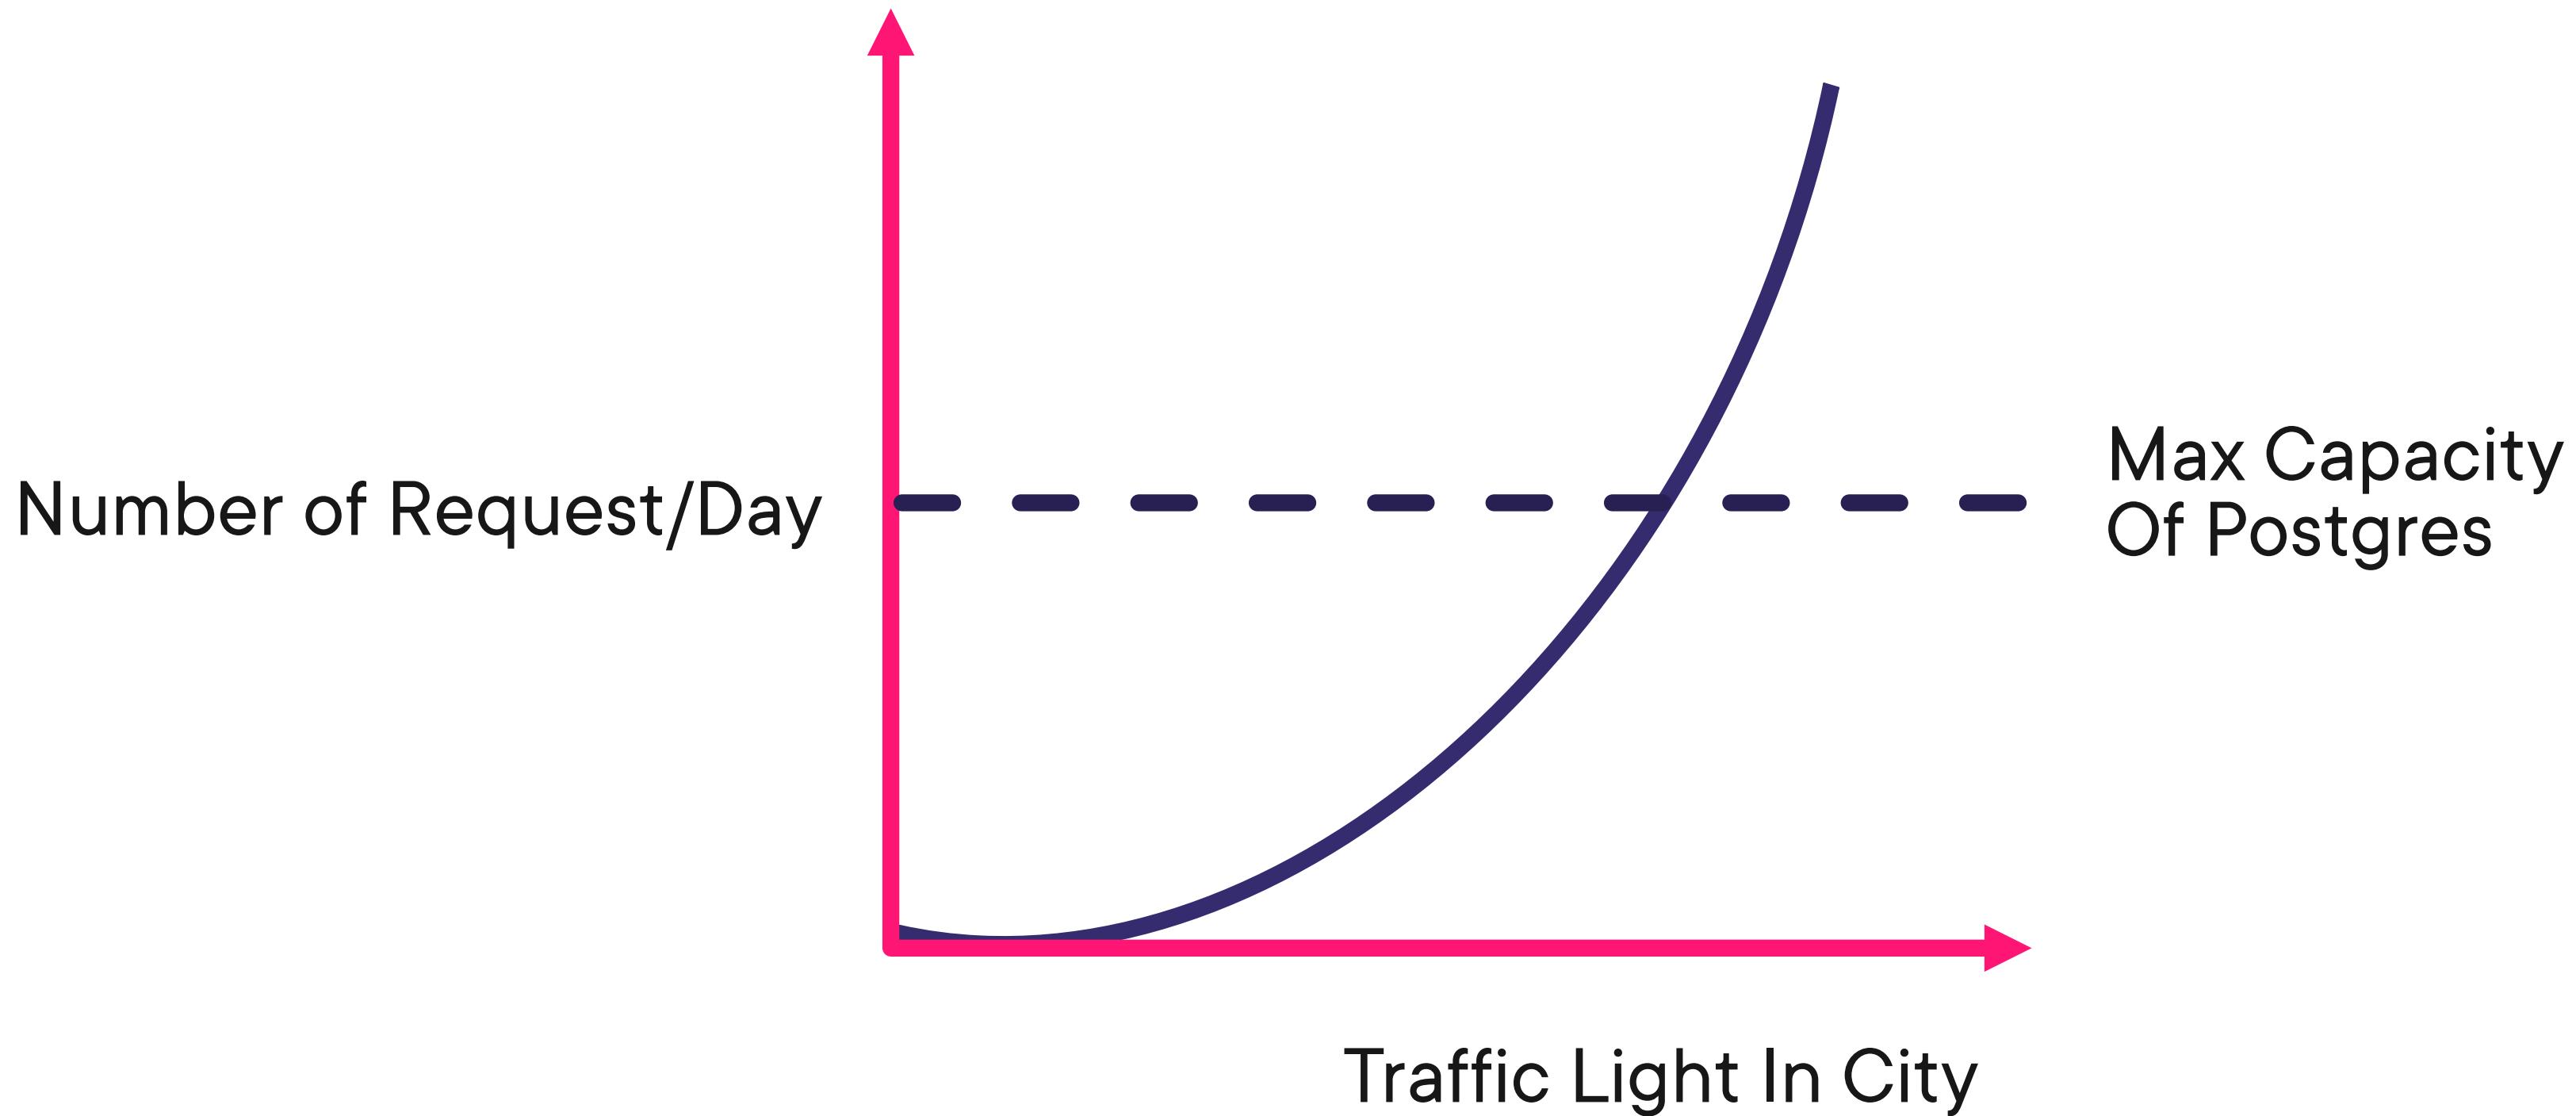
Why Kafka?



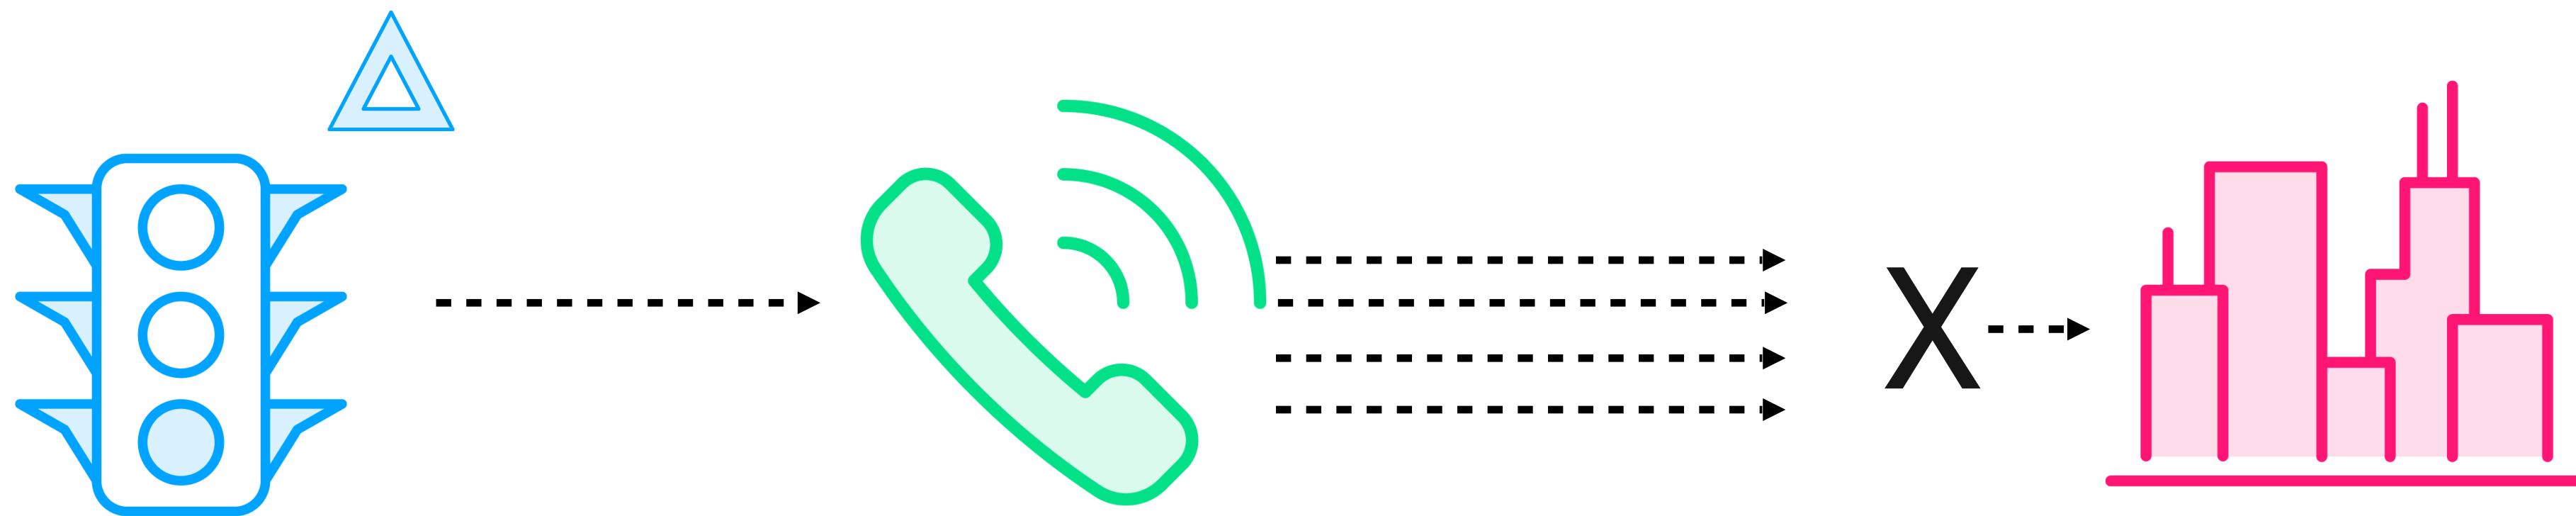
A Case Study: Traffic Lights



A Case Study: Traffic Lights



A Case Study: Traffic Lights

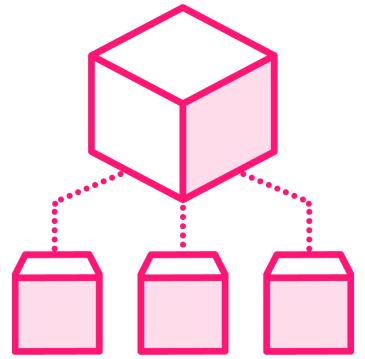


A Case Study: 911 Systems

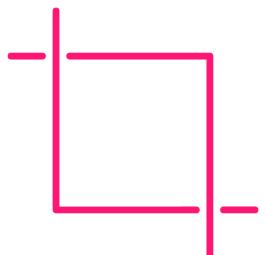
Kafka comes to solve this!



What Is Kafka?



Event Streaming Platform



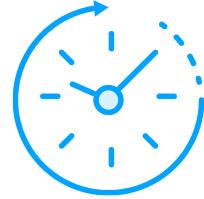
Scalable



High Throughput



Difference between Kafka and the Rest of Solutions



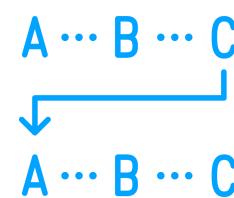
RabbitMQ: As Kafka uses topics in an unidirectional log without handshake, the **difference in performance** is brutal, easily 1000x



Mulesoft: Kafka **makes it simple** to create a channel of communication simply by producing a message the topic will be autocreated



Redis: Redis **does not handle streaming data** nor has a history past messages



Amazon SQS: It based on queues, way better tan RabbitMQ but still there is an **order of magnitude in throughput** and speed between Kafka and SQS. Plus SQS is not exactly cheap



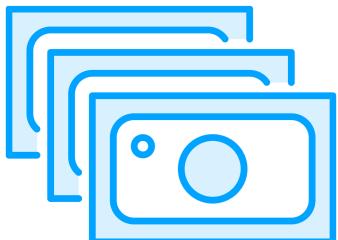
Difference between Kafka and the Rest of Solutions



Amazon SNS: Very, very similar; but **way pricer**



Azure Streaming Services: This service in particular is **not exactly meant for pub-sub in general**



Google Pub/Sub: The same as Amazon SNS, very similar but **pricer**

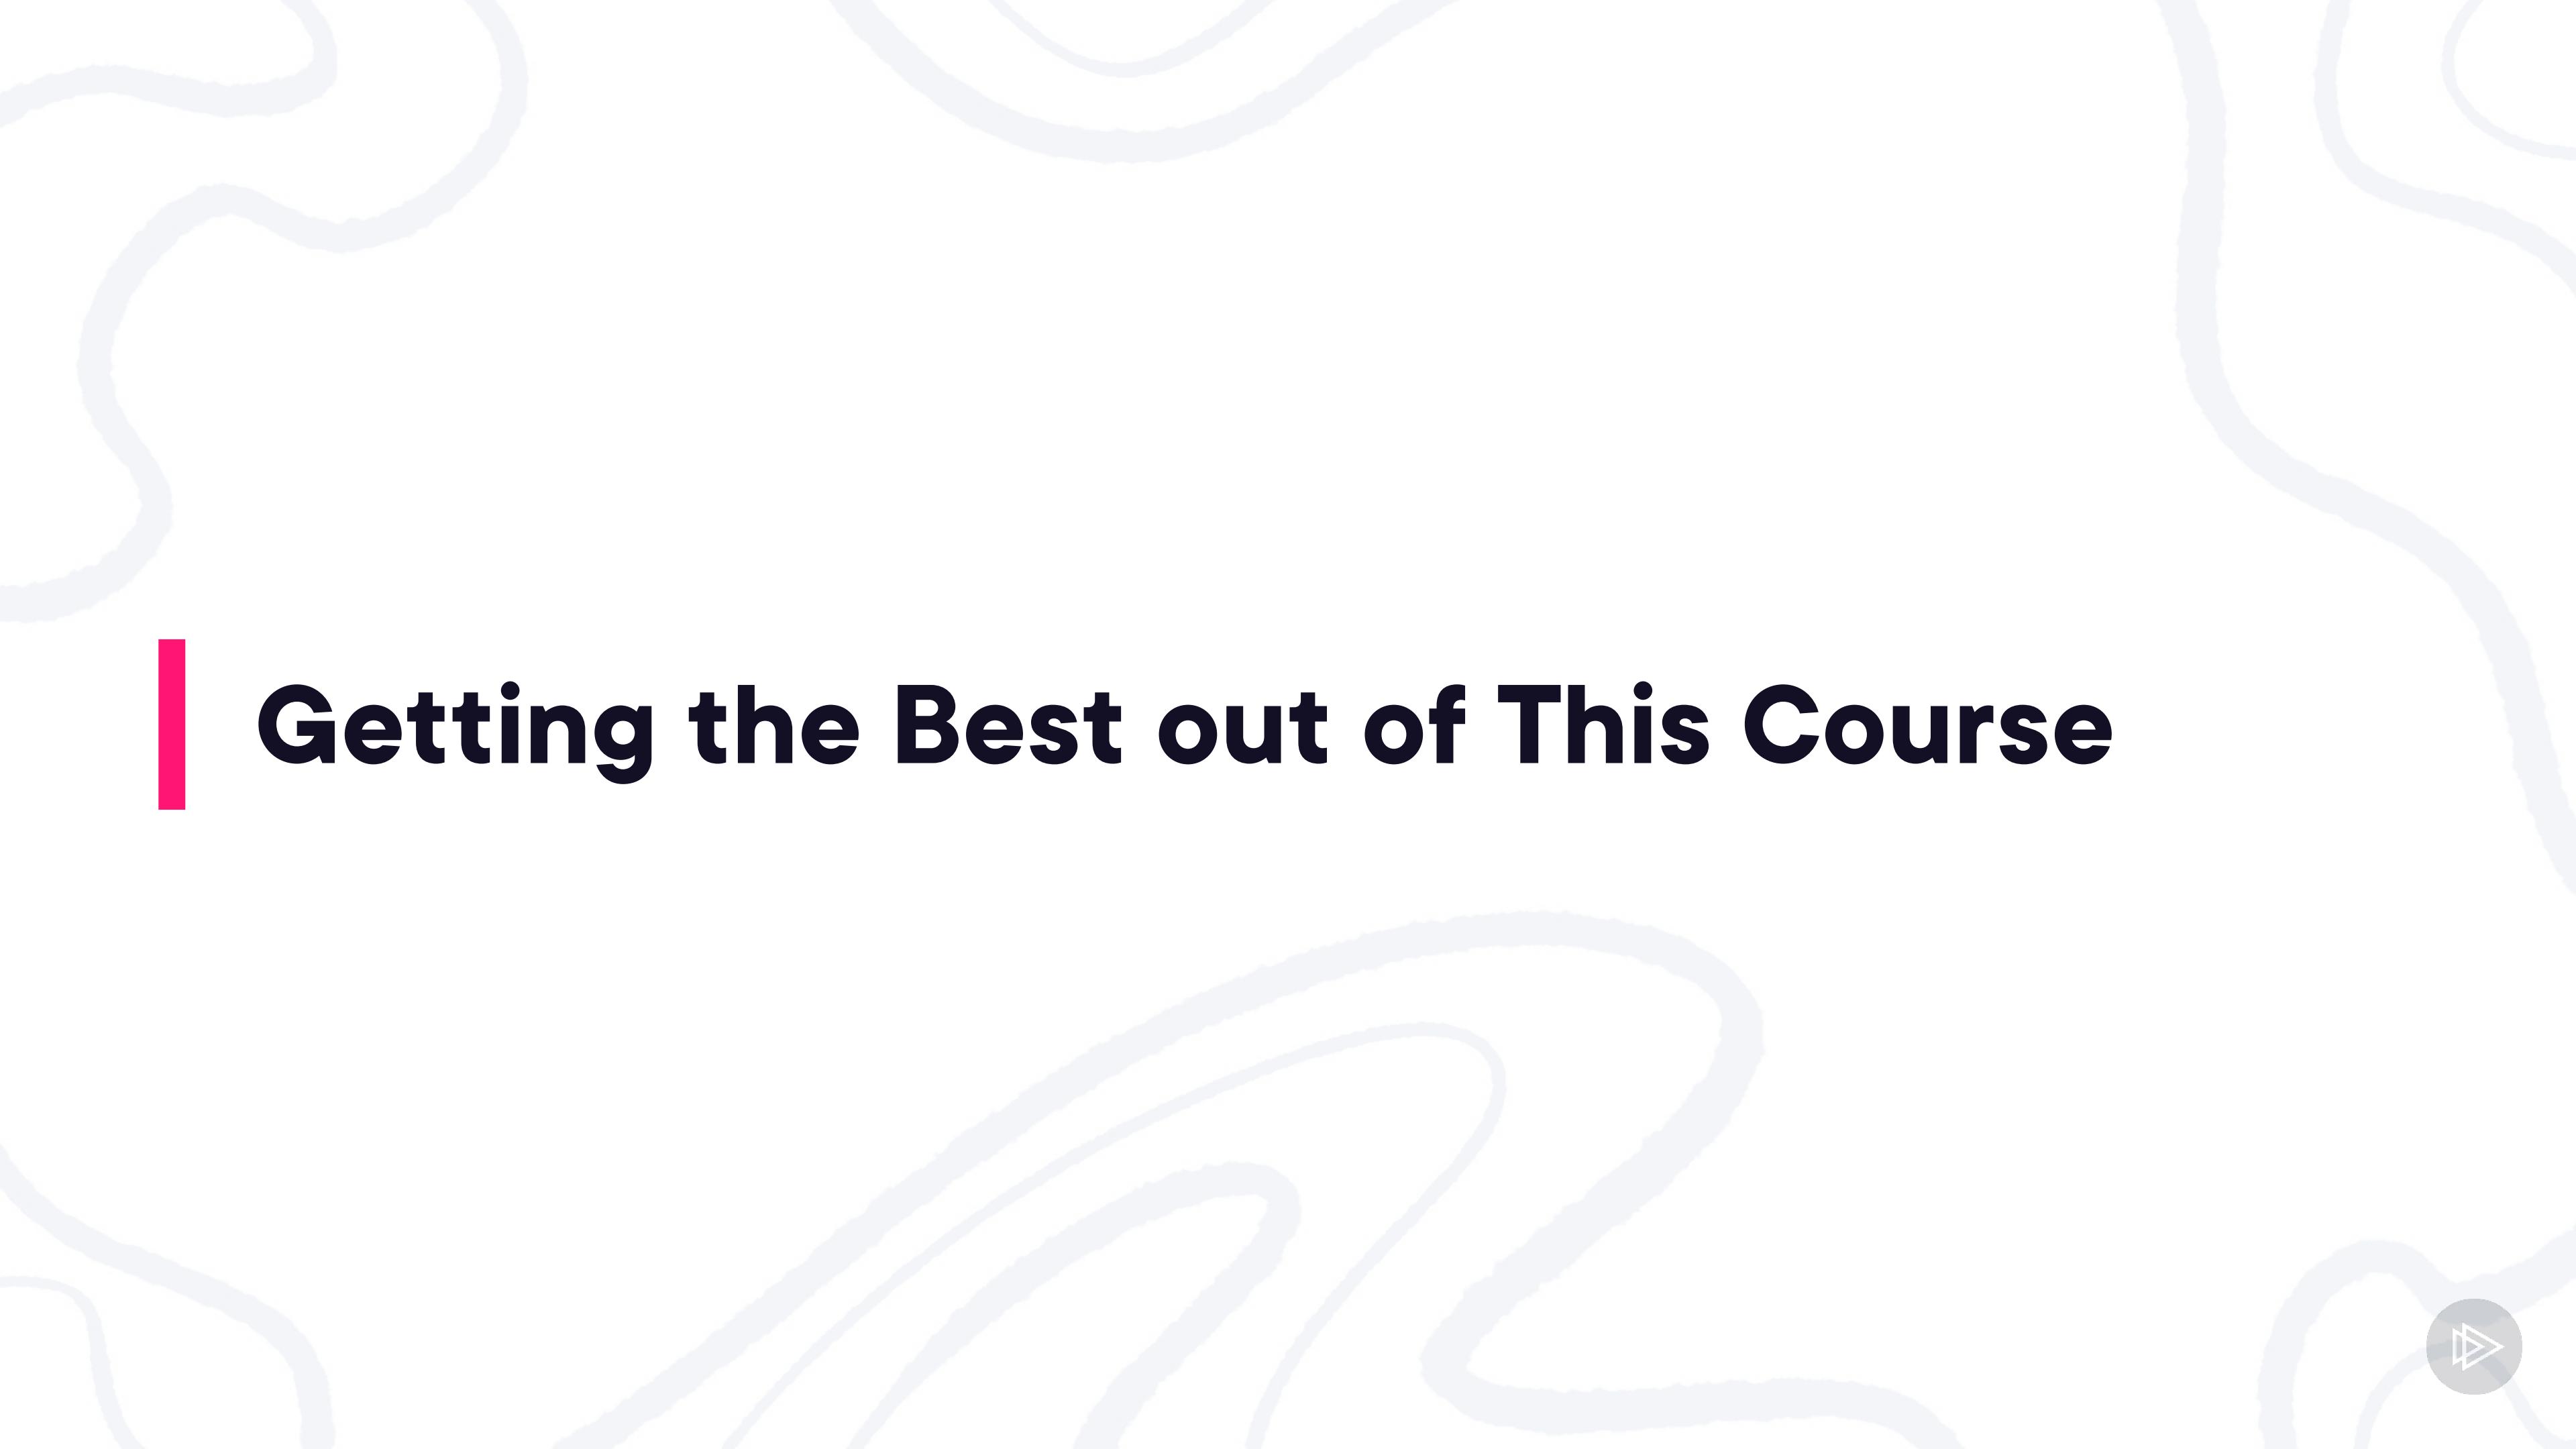


Demo: Introducing Globomantics Case Study



- Amazon bill went up 300x
- Apache Kafka can replace such services
- It can reduce the bill significantly





Getting the Best out of This Course



Prerequisites

Prerequisites

- Basic Java
- Basic Docker

Course Suggestions

- Java SE 17 Fundamentals
- Docker: The Big Picture



More Resources

[Table of contents](#)[Description](#)[Transcript](#)[Exercise files](#)[**Discussion**](#)[Learning Check](#)[Related Courses](#)

AUTHOR TOOLS

 Home Analytics Author's nest Author kitAxel Sirota 

Pluralsight Author

 Following

133 Followers

Ask me on Twitter!

@AxelSirota



GitHub Repository for the Course

	axel-sirota	add all	429f0d2	11 hours ago	 4 commits
	module3	add all		11 hours ago	
	module4	add all		11 hours ago	
	.gitignore	module 3 demos		3 days ago	
	LICENSE	Initial commit		4 days ago	
	README.md	Initial commit		4 days ago	
	requirements.txt	add structure		2 days ago	
	workspace.yml	module 3 demos		3 days ago	

<https://github.com/axel-sirota/getting-started-kafka>



Version Check



Version Check



**Apache Kafka – Confluent Version 7.4
(equivalent to Kafka open source 3.4)**

Apache Zookeeper 3.8.1

Kafka Rest Proxy – Confluent Version 7.4

Kafka Connect – Confluent Version 7.4

ksqldb 0.29

Schema Registry – Confluent Version 7.4

Kafka Streams – Confluent Version 7.4



Version Check



This course is 100% applicable to:

- Apache Kafka – Confluent Version ≥ 6.0
(equivalent to Kafka open source version ≥ 2.6)
- Apache Zookeeper 3.4.10 through 3.5.9 (3.5.9 recommended)
- Kafka Rest Proxy – Confluent Version ≥ 6.0
- Kafka Connect – Confluent Version ≥ 6.0
- ksqlDB 0.2
- Schema Registry – Confluent Version ≥ 6.0
- Kafka Streams – Confluent Version ≥ 6.0

<https://docs.confluent.io/platform/current/installation/versions-interoperability.html>



Version Check



This course is NOT applicable to:

- Apache Kafka - Confluent Version < 6.0
- Apache Zookeeper < 3.4.10

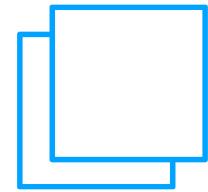




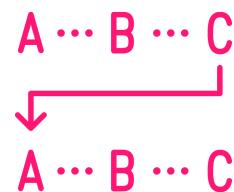
What Will You Be Able to Do When Finishing This Course?



After This Course, Not Only:



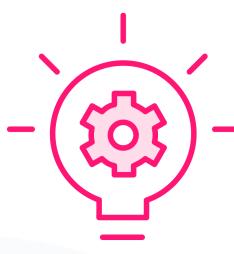
You will be able to create applications that produces and consumes messages from a topic in Kafka



Actually you will be fully equipped to create an event driven application with full governance over how messages flow



You will be able to expose this messages as an API to perform queries over it, or even create tables to understand how Streams of data behave



Finally, you will also have the knowledge to secure your cluster and scale it up to have high availability



Outline Of The Course



Kafka Architecture

Kafka Streams

Perform administrative tasks

Creating and Configuring Producers and Consumers

Kafka Connect

**Final Thoughts
Key Takeaways**

